

PTHREAD PROGRAMMING

What is Threading?

- Threading is the creation and management of multiple units of execution within a single process.
- Threading is a significant source of programming error, through the introduction of data races and deadlocks

Binaries, Processes, and Threads

- **Binaries** are dormant programs residing on a storage medium, ready to execute but not yet in motion.
- **Processes** are the operating system abstraction representing those binaries in action.
- **Threads** are the unit of execution within a process.
 - If a process contains but one thread, there is only a single unit of execution in the process. We call such processes **single threaded**.
 - If a process contains more than one thread, then there is more than one thing going on at once. We call such processes **multithreaded**.

Two Fundamental Virtualized Abstractions

- virtual memory and
- a virtualized processor
- Together, they give the illusion to each running process that it alone consumes the machine's resources.
 - Virtualized memory is associated with the process and not the thread.
 - Conversely, a virtualized processor is associated with threads and not processes.

- A thread consists of the information necessary to represent an execution context within a process.
- This includes a thread ID that identifies the thread within a process, a set of register values, a stack, a scheduling priority and policy, a signal mask, an errno variable, and thread-specific data.
- The threads interfaces we're about to see are from POSIX.1-2001. The threads interfaces, also known as **"pthreads"** for **"POSIX threads,"**

Thread Identification

- Just as every process has a process ID, every thread has a thread ID.
- Unlike the process ID, which is unique in the system, the thread ID has significance only within the context of the process to which it belongs.
- Recall that a process ID, represented by the `pid_t` data type, is a non-negative integer.
- A thread ID is represented by the **`pthread_t`** data type.

The Functions

- Checking equality

```
#include <pthread.h>
int pthread_equal(pthread_t tid1, pthread_t tid2);
Returns: nonzero if equal, 0 otherwise
```

- A thread can obtain its own thread ID by calling the pthread_self function.

```
#include <pthread.h>
pthread_t pthread_self(void);
Returns: the thread ID of the calling thread
```

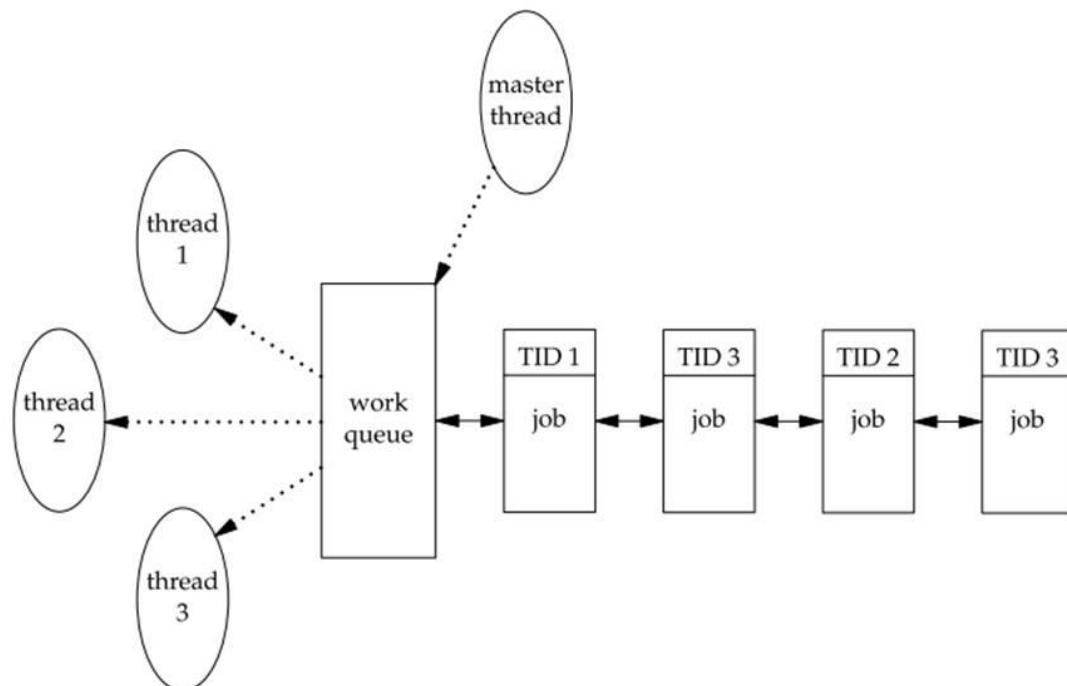
Thread Creation

- threads can be created by calling the pthread_create function.

```
#include <pthread.h>
int pthread_create(pthread_t *restrict tidp,
const pthread_attr_t *restrict attr, void *(* start_rtn )(void *), void *restrict arg );
Returns: 0 if OK, error number on failure
```


Attributes

- The memory location pointed to by *tidp* is set to the thread ID of the newly created thread when `pthread_create` returns successfully.
- The *attr* argument is used to customize various thread attributes.
- The newly created thread starts running at the address of the *start_rtn* function.
- This function takes a single argument, *arg*, which is a typeless pointer.



Example

```
#include <pthread.h>

pthread_t ntid;
void printids(const char *s)
{
    pid_t    pid;
    pthread_t tid;
    pid = getpid();
    tid = pthread_self();
    printf("%s pid %lu tid %lu (0x%lx)\n", s, (unsigned long)pid,
        (unsigned long)tid, (unsigned long)tid);
}

void * thr_fn(void *arg)
{
    printids("new thread: ");
}

int main(void)
{
    int    err;
    err = pthread_create(&ntid, NULL, thr_fn,
        NULL);
    if (err != 0)
        err_exit(err, "can't create thread");
    printids("main thread:");
    sleep(1);
    exit(0);
}
```

Thread Termination

- If any thread within a process calls `exit`, `_Exit`, or `_exit`, then the entire process terminates.
- A single thread can exit in three ways, thereby stopping its flow of control, without terminating the entire process.
 - The thread can simply return from the start routine. The return value is the thread's exit code.
 - The thread can be canceled by another thread in the same process.

The `pthread_exit()`

```
#include <pthread.h>
void pthread_exit(void * rval_ptr);
```

- The `rval_ptr` argument is a typeless pointer, similar to the single argument passed to the start routine.

Joining Threads

- This pointer is available to other threads in the process by calling the
- `pthread_join` function.

```
#include <pthread.h>
```

```
int pthread_join(pthread_t thread, void ** rval_ptr);
```

Returns: 0 if OK, error number on failure

- The calling thread will block until the specified thread calls `pthread_exit`, returns from its start routine, or is canceled.
- If the thread simply returned from its start routine, `rval_ptr` will contain the return code.
- If the thread was canceled, the memory location specified

Demonstrating Exit code

```
#include <pthread.h>
```

```
void * thr_fn1(void *arg)
```

```
{  
    printf("thread 1 returning\n");  
    return((void *)1);  
}
```

```
void * thr_fn2(void *arg)
```

```
{  
    printf("thread 2 exiting\n");  
    pthread_exit((void *)2);  
}
```



```

int main(void)
{
    int err;
    pthread_t tid1, tid2;
    void *tret;
    err = pthread_create(&tid1, NULL, thr_fn1, NULL);
    if (err != 0)
        printf("can't create thread 1");
    err = pthread_create(&tid2, NULL, thr_fn2, NULL);
    if (err != 0)
        printf("can't create thread 2");
    err = pthread_join(tid1, &tret);
    if (err != 0)
        printf("can't join with thread 1");
    printf("thread 1 exit code %ld\n", (long)tret);
    err = pthread_join(tid2, &tret);
    if (err != 0)
        printf("can't join with thread 2");
    printf("thread 2 exit code %ld\n", (long)tret);
    exit(0);
}

```

Passing more than one value

- The typeless pointer passed to `pthread_create` and `pthread_exit` can be used to pass more than a single value.
- The pointer can be used to pass the address of a structure containing more complex information.
- Be careful that the memory used for the structure is still valid when the caller has completed.

Example

```
#include <pthread.h>
struct foo {
    int a, b, c, d;
};
Void printfoo(const char *s, const struct foo *fp)
{
    printf("%s", s);
    printf(" structure at 0x%lx\n", (unsigned long)fp);
    printf(" foo.a = %d\n", fp->a);
    printf(" foo.b = %d\n", fp->b);
    printf(" foo.c = %d\n", fp->c);
    printf(" foo.d = %d\n", fp->d);
}
void * thr_fn1(void *arg)
{
    struct foo foo = {1, 2, 3, 4};
    printfoo("thread 1:\n", &foo);
    pthread_exit((void *)&foo);
}
void * thr_fn2(void *arg)
{
    int main(void)
    {
        int err;
        pthread_t tid1, tid2;
        struct foo *fp;
        err = pthread_create(&tid1, NULL, thr_fn1, NULL);
        if (err != 0)
            printf("can't create thread 1");
        err = pthread_join(tid1, (void *)&fp);
        if (err != 0)
            printf("can't join with thread 1");
        sleep(1);
        printf("parent starting second thread\n");
        err = pthread_create(&tid2, NULL, thr_fn2, NULL);
        if (err != 0)
            printf("can't create thread 2");
        sleep(1);
        printfoo("parent:\n", fp);
    }
}
```

Cancelling Threads

- One thread can request that another in the same process be canceled by calling the `pthread_cancel` function.

```
#include <pthread.h>
int pthread_cancel(pthread_t tid);
```

Returns: 0 if OK, error number on failure

Thread cleanup handlers.

- A thread can arrange for functions to be called when it exits, similar to the way that the `atexit` function can be used by a process to arrange that functions are to be called when the process exits.
- The functions are known as *thread cleanup handlers*.
- More than one cleanup handler can be established for a thread.
- *The handlers are recorded in a stack*, which


```
#include <pthread.h>
void pthread_cleanup_push(void (*rtn)(void *), void * arg);
void pthread_cleanup_pop(int execute);
```

– The `pthread_cleanup_push` function schedules the cleanup function, `rtn`, to be called with the single argument, `arg`, when the thread performs one of the following actions:

- Makes a call to `pthread_exit`
- Responds to a cancellation request
- Makes a call to `pthread_cleanup_pop` with a nonzero `execute` argument

Similarities between thread and Process functions

Process primitive	Thread primitive	Description
<code>fork</code>	<code>pthread_create</code>	create a new flow of control
<code>exit</code>	<code>pthread_exit</code>	exit from an existing flow of control
<code>waitpid</code>	<code>pthread_join</code>	get exit status from flow of control
<code>atexit</code>	<code>pthread_cleanup_push</code>	register function to be called at exit from flow of control
<code>getpid</code>	<code>pthread_self</code>	get ID for flow of control
<code>abort</code>	<code>pthread_cancel</code>	request abnormal termination of flow of control